

Final Submission - Music Recommendation Systems

Executive Summary

- What are the key takeaways?
- What are the key next steps?

Problem and Solution Summary

- What problem was being solved?
- What are the key points that describe the final proposed solution design?
- Why is this a 'valid' solution that is likely to solve the problem?

Recommendations for Implementation

- What are some key recommendations to implement the solutions?
- What are the key actionables for stakeholders?
- What is the expected benefit and/or costs? What are the key risks and challenges?
- What further analysis needs to be done or what other associated problems need to be solved?

Milestone 1

Data Dictionary

The core data is the Taste Profile Subset released by the Echo Nest as part of the Million Song Dataset. There are two files in this dataset. The first file contains the details about the song id, titles, release, artist name, and the year of release. The second file contains the user id, song id, and the play count of users.

song_data

song_id - A unique id given to every song

title - Title of the song

Release - Name of the released album

Artist_name - Name of the artist

year - Year of release

count_data

user_id - A unique id given to the user

song_id - A unique id given to the song

play_count - Number of times the song was played

Data Source

<http://millionsongdataset.com/>

Important Notes

- This notebook can be considered a guide to refer to while solving the problem. The evaluation will be as per the Rubric shared for each Milestone. Unlike previous courses, it does not follow the pattern of the graded questions in different sections. This notebook would give you a direction on what steps need to be taken to get a feasible solution to the problem. Please note that this is just one way of doing this. There can be other 'creative' ways to solve the problem, and we encourage you to feel free and explore them as an 'optional' exercise.
- In the notebook, there are markdown cells called Observations and Insights. It is a good practice to provide observations and extract insights from the outputs.
- The naming convention for different variables can vary. Please consider the code provided in this notebook as a sample code.
- All the outputs in the notebook are just for reference and can be different if you follow a different approach.
- There are sections called Think About It in the notebook that will help you get a better understanding of the reasoning behind a particular technique/step. Interested learners can take alternative approaches if they want to explore different techniques.

Importing Libraries and the Dataset

```
In [1]: from google.colab import drive  
drive.mount('/content/drive')
```

Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount("/content/drive", force_remount=True).

```
In [2]: import warnings  
warnings.filterwarnings('ignore')  
  
import numpy as np  
import pandas as pd  
  
import matplotlib.pyplot as plt  
  
import seaborn as sns  
  
from sklearn.metrics.pairwise import cosine_similarity
```

```
from collections import defaultdict
from sklearn.metrics import mean_squared_error
```

Load the dataset

```
In [3]: count_df = pd.read_csv('/content/drive/MyDrive/count_data.csv')
song_df = pd.read_csv('/content/drive/MyDrive/song_data.csv')
```

Understanding the data by viewing a few observations

```
In [4]: count_df.head(10)
```

```
Out[4]:   Unnamed: 0          user_id        song_id  play_count
0         0  b80344d063b5ccb3212f76538f3d9e43d87dca9e  SOAKIMP12A8C130995      1
1         1  b80344d063b5ccb3212f76538f3d9e43d87dca9e  SOBBMDR12A8C13253B      2
2         2  b80344d063b5ccb3212f76538f3d9e43d87dca9e  SOBXHDL12A81C204C0      1
3         3  b80344d063b5ccb3212f76538f3d9e43d87dca9e  SOBYHAJ12A6701BF1D      1
4         4  b80344d063b5ccb3212f76538f3d9e43d87dca9e  SODACBL12A8C13C273      1
5         5  b80344d063b5ccb3212f76538f3d9e43d87dca9e  SODDNQT12A6D4F5F7E      5
6         6  b80344d063b5ccb3212f76538f3d9e43d87dca9e  SODXRTY12AB0180F3B      1
7         7  b80344d063b5ccb3212f76538f3d9e43d87dca9e  SOFGUAY12AB017B0A8      1
8         8  b80344d063b5ccb3212f76538f3d9e43d87dca9e  SOFRQTD12A81C233C0      1
9         9  b80344d063b5ccb3212f76538f3d9e43d87dca9e  SOHQWYZ12A6D4FA701      1
```

```
In [5]: song_df.head(10)
```

```
Out[5]:       song_id           title           release  artist_name  year
0  SOQMMHC12AB0180CB8  Silent Night  Monster Ballads X-Mas  Faster Pussy cat  2003
1  SOVFVAK12A8C1350D9  Tanssi vaan  Karkuteillä  Karkkiautomaatti  1995
2  SOGTUKN12AB017F4F1  No One Could Ever  Butter  Hudson Mohawke  2006
3  SOBNYVR12A8C13558C  Si Vos Querés  De Culo  Yerba Brava  2003
4  SOHSBXH12A8C13B0DF  Tangle Of Aspens  Rene Ablaze Presents Winter Sessions  Der Mystic  0
5  SOZVAPQ12A8C13B63C  Symphony No. 1 G minor "Sinfonie Serieuse"/All...  Berwald: Symphonies Nos. 1/2/3/4  David Montgomery  0
6  SOQVRHI12A6D4FB2D7  We Have Got Love  Strictly The Best Vol. 34  Sasha / Turbulence  0
7  SOEYRFT12AB018936C  2 Da Beat Ch'yall  Da Bomb  Kris Kross  1993
8  SOPMIYT12A6D4F851E  Goodbye  Danny Boy  Joseph Locke  0
9  SOJCFMH12A8C13B0C2  Mama_ mama can't you see ?  March to cadence with the US marines  The Sun Harbor's Chorus-Documentary Recordings  0
```

Let us check the data types and missing values of each column

```
In [6]: count_df['user_id'].nunique()
```

```
Out[6]:76353
```

```
In [7]: song_df['song_id'].nunique()
```

```
Out[7]:999056
```

```
In [8]: count_df['song_id'].nunique()
```

```
Out[8]:10000
```

```
In [9]: count_df['play_count'].nunique()
```

```
Out[9]:295
```

```
In [10]: song_df['release'].nunique()
```

```
Out[10]:149288
```

```
In [11]: song_df['year'].nunique()
```

```
Out[11]:90
```

```
In [12]: song_df['artist_name'].nunique()
```

```
Out[12]:72665
```

```
In [13]: song_df.describe()
```

```
Out[13]:          year
count  1000000.000000
```

```
mean    1030.325652
std     998.745002
min     0.000000
25%    0.000000
50%    1969.000000
75%    2002.000000
max     2011.000000
```

```
In [14]: count_df.describe()
```

```
Out[14]:      Unnamed: 0      play_count
count  2.000000e+06  2.000000e+06
mean   9.999995e+05  3.045485e+00
std    5.773504e+05  6.579720e+00
min    0.000000e+00  1.000000e+00
25%   4.999998e+05  1.000000e+00
50%   9.999995e+05  1.000000e+00
75%   1.499999e+06  3.000000e+00
max   1.999999e+06  2.213000e+03
```

```
In [15]: song_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1000000 entries, 0 to 999999
Data columns (total 5 columns):
 #   Column      Non-Null Count  Dtype  
 ---  --          -----          ----  
 0   song_id     1000000 non-null   object 
 1   title       999985 non-null   object 
 2   release     999995 non-null   object 
 3   artist_name 1000000 non-null   object 
 4   year        1000000 non-null   int64  
dtypes: int64(1), object(4)
memory usage: 38.1+ MB
```

```
In [16]: song_df.isnull().sum()
```

```
Out[16]:song_id      0
         title       15
         release     5
         artist_name 0
         year        0
         dtype: int64
```

```
In [17]: count_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2000000 entries, 0 to 1999999
Data columns (total 4 columns):
 #   Column      Dtype  
 ---  --          ----  
 0   Unnamed: 0   int64 
 1   user_id     object 
 2   song_id     object 
 3   play_count  int64  
dtypes: int64(2), object(2)
memory usage: 61.0+ MB
```

```
In [18]: count_df.isnull().sum()
```

```
Out[18]:Unnamed: 0      0
         user_id     0
         song_id     0
         play_count  0
         dtype: int64
```

Observations and Insights:

- Overall there are 1,000,000 songs in the file and 2,000,000 interactions in the count file. In the song file there are a total 20 of missing values with 15 coming from the song titles column and the other 5 that derive from the album release column. Most of the data types in the two files are objective with the other few being integer data types, including how many times a song was played.
- One thing I noticed is that from the first 10 records in the song file is that there are several 0's under the year column, even though they are zero missing values when looking at the data types. This could possibly impact the users recommendations if the time period is a key aspect of what a user listens to. For example, if a user listens to a lot of music from the 1970's, the best recommendations from that era may not be provided to the user because that information is missing.
- From the count file there doesn't seem to be anything that stands out like missing values.

```
In [19]: df = pd.merge(count_df, song_df.drop_duplicates(['song_id']), on="song_id", how="left")
In [20]: df.drop(['Unnamed: 0'], axis=1, inplace=True)
In [21]: df.info()

<class 'pandas.core.frame.DataFrame'>
Int64Index: 2000000 entries, 0 to 1999999
Data columns (total 7 columns):
 #   Column      Dtype  
--- 
 0   user_id     object 
 1   song_id     object 
 2   play_count  int64  
 3   title       object 
 4   release     object 
 5   artist_name object 
 6   year        int64  
dtypes: int64(2), object(5)
memory usage: 122.1+ MB
In [22]: df.head()
```

```
Out[22]:
```

	user_id	song_id	play_count	title	release	artist_name	year
0	b80344d063b5ccb3212f76538f3d9e43d87dca9e	SOAKIMP12A8C130995	1	The Cove	Thicker Than Water	Jack Johnson	0
1	b80344d063b5ccb3212f76538f3d9e43d87dca9e	SOBBMDR12A8C13253B	2	Entre Dos Aguas	Flamenco Para Niños	Paco De Lucia	1976
2	b80344d063b5ccb3212f76538f3d9e43d87dca9e	SOBXHDL12A81C204C0	1	Stronger	Graduation	Kanye West	2007
3	b80344d063b5ccb3212f76538f3d9e43d87dca9e	SOBYHAJ12A6701BF1D	1	Constellations	In Between Dreams	Jack Johnson	2005
4	b80344d063b5ccb3212f76538f3d9e43d87dca9e	SODACBL12A8C13C273	1	Learn To Fly	There Is Nothing Left To Lose	Foo Fighters	1999

```
In [23]: df['song_id'].nunique()
```

```
Out[23]: 10000
```

```
In [24]: df['user_id'].nunique()
```

```
Out[24]: 76353
```

```
In [25]: df['artist_name'].nunique()
```

```
Out[25]: 3375
```

Think About It: As the user_id and song_id are encrypted. Can they be encoded to numeric features?

- Encrypting the user_id and song_id columns in the files I believe can be encoded to numeric features and would be easier to visualize. Their original objective data type has each entry as being a large group of numbers and letters, so encoding them to just being a numeric feature will be easier to understand for future visualization and programming.

```
In [26]: !pip install scipy
```

```
Looking in indexes: https://pypi.org/simple, https://us-python.pkg.dev/colab-wheels/public/simple/
Requirement already satisfied: scipy in /usr/local/lib/python3.8/dist-packages (1.7.3)
Requirement already satisfied: numpy<1.23.0,>=1.16.5 in /usr/local/lib/python3.8/dist-packages (from scipy) (1.21.6)
In [27]: from pandas.core.arrays.interval import le
         from sklearn.preprocessing import LabelEncoder
         from sklearn.metrics.pairwise import cosine_similarity
```

```
le = LabelEncoder()

df['user_id'] = le.fit_transform(df['user_id'])

df['song_id'] = le.fit_transform(df['song_id'])
```

```
In [28]: df.head()
```

```
Out[28]:
```

	user_id	song_id	play_count	title	release	artist_name	year
0	54961	153	1	The Cove	Thicker Than Water	Jack Johnson	0
1	54961	413	2	Entre Dos Aguas	Flamenco Para Niños	Paco De Lucia	1976
2	54961	736	1	Stronger	Graduation	Kanye West	2007
3	54961	750	1	Constellations	In Between Dreams	Jack Johnson	2005
4	54961	1188	1	Learn To Fly	There Is Nothing Left To Lose	Foo Fighters	1999

Think About It: As the data also contains users who have listened to very few songs and vice versa, is it required to filter the data so that it contains users who have listened to a good count of songs and vice versa?

- I think collectively it makes sense to include both types of users in terms of how many songs they have listened to. Regardless of how many listens a user has I think there is enough information there to include both types of users where a reliable recommendation system can be built. I think even if a user has only listened to a small amount of songs that there could still be some commonalities between the listened

songs where a recommendation system can be built and wouldn't be too much different from a user who has listened to a lot of songs. Despite this, more listens would create a better recommendation system with the amount of data to use.

In [29]: users = df.user_id

```
ratings_count = dict()

for user in users:

    if user in ratings_count:
        ratings_count[user] += 1

    else:
        ratings_count[user] = 1
```

In [30]: RATINGS_CUTOFF = 90

```
remove_users = []
for user, num_ratings in ratings_count.items():
    if num_ratings < RATINGS_CUTOFF:
        remove_users.append(user)
df = df.loc[~df.user_id.isin(remove_users)]
```

In [31]: songs = df.song_id

```
ratings_count = dict()

for song in songs:

    if song in ratings_count:
        ratings_count[song] += 1

    else:
        ratings_count[song] = 1
```

In [32]: RATINGS_CUTOFF = 120

```
remove_songs = []
for song, num_ratings in ratings_count.items():
    if num_ratings < RATINGS_CUTOFF:
        remove_songs.append(song)

df_final = df.loc[ ~ df.song_id.isin(remove_songs)]
```

In [33]: df_final=df_final[df_final.play_count<=5]

In [34]: df_final.info()

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 117876 entries, 200 to 1999765
Data columns (total 7 columns):
 #   Column      Non-Null Count  Dtype  
 ---  --          --          --      
 0   user_id     117876 non-null   int64  
 1   song_id     117876 non-null   int64  
 2   play_count  117876 non-null   int64  
 3   title       117876 non-null   object 
 4   release     117876 non-null   object 
 5   artist_name 117876 non-null   object 
 6   year        117876 non-null   int64  
dtypes: int64(4), object(3)
memory usage: 7.2+ MB
```

In [35]: df_final.shape

Out[35]:(117876, 7)

In [36]: df_final.size

Out[36]:825132

In [37]: df_final.head()

```
Out[37]:    user_id  song_id  play_count           title      release  artist_name  year
        200      6958      447         1  Daisy And Prudence  Distillation  Erin McKeown  2000
        202      6958      512         1  The Ballad of Michael Valentine  Sawdust      The Killers  2004
        203      6958      549         1  I Stand Corrected (Album)  Vampire Weekend  Vampire Weekend  2007
        204      6958      703         1  They Might Follow You  Tiny Vipers      Tiny Vipers  2007
        205      6958      719         1  Monkey Man  You Know I'm No Good  Amy Winehouse  2007
```

Exploratory Data Analysis

Let's check the total number of unique users, songs, artists in the data

Total number of unique user id

```
In [38]: print('Number of unique USERS in Raw data = ', df_final['user_id'].nunique())
```

Number of unique USERS in Raw data = 3155

Total number of unique song id

```
In [39]: print('Number of unique SONGS in Raw data = ', df_final['song_id'].nunique())
```

Number of unique SONGS in Raw data = 563

Total number of unique artists

```
In [40]: print('Number of unique ARTISTS in Raw data = ', df_final['artist_name'].nunique())
```

Number of unique ARTISTS in Raw data = 232

Observations and Insights:

- There is a decent amount of unique users that have listened to at least 90 songs in the dataset but after the reduction of users from the minimum requirement of songs listened to I am surprised it is still kind of high. My guess is that many users just listen to music casually without repeating songs or the recommendation system in place from past interactions isn't entirely suited to their music interests.
- There is a much smaller number of unique songs in the revised dataset which indicates that they are getting a lot of users listening to them. Maybe these are attractive songs for the recommendation system to suggest in terms of past interactions but not popular songs in the mainstream where they will receive more than one listen on a streaming service.
- The total number of unique artists is even smaller than the unique songs so based off the revised dataset maybe these are lesser known artists in terms of popularity or its possible they don't have much music available in the dataset. Maybe they also did not receive enough listens of their music so they did not show up in the condensed dataset.

Let's find out about the most interacted songs and interacted users

Most interacted songs

```
In [41]: most_rated = df_final.groupby('song_id').size().sort_values(ascending = False)[:10]
most_rated
```

Out[41]: song_id

8582	751
352	748
2220	713
1118	662
4152	652
5531	618
4448	609
6189	606
6293	583
1334	570

dtype: int64

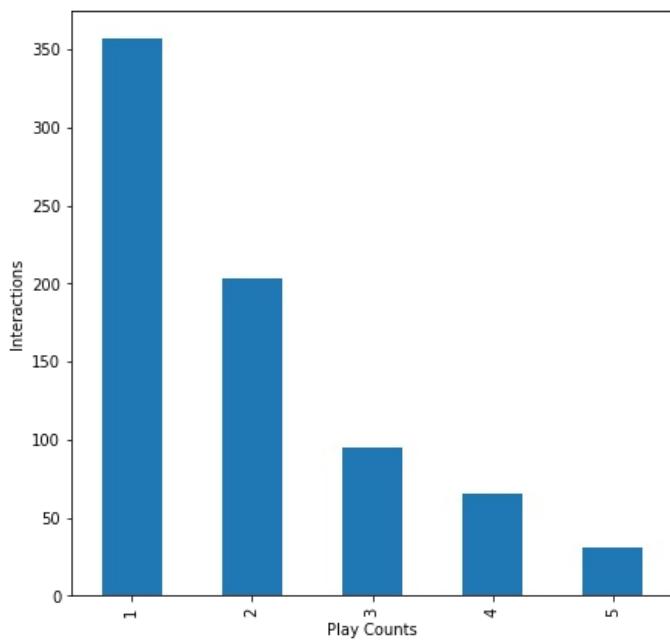
```
In [42]: # Plotting distributions of play counts for 751 interactions with song_id 8582
```

```
plt.figure(figsize = (7, 7))

df_final[df_final['song_id'] == 8582]['play_count'].value_counts().plot(kind = 'bar')

plt.xlabel('Play Counts')
plt.ylabel('Interactions')

plt.show()
```



- Plotting the song with the most ratings (song_id 8582 with 751 interactions) shows an interesting breakdown as to how many times the song was listened to with the cap of 5 interactions maximum. Almost half of the interactions had just one listen and decreased with every possible interaction. So overall, despite the song being the most interacted with, it is not a song that streamers will likely go back and listen to multiple

times.

Most interacted users

```
In [43]: most_rated = df_final.groupby('user_id').size().sort_values(ascending = False)[:10]
most_rated
```

```
Out[43]:user_id
```

61472	243
15733	227
37049	202
9570	184
23337	177
10763	176
9097	175
26616	175
43041	174
65994	171

dtype: int64

Interacted songs by year

```
In [44]: most_rated = df_final.groupby('year').size().sort_values(ascending = False)[:10]
most_rated
```

```
Out[44]:year
```

0	20649
2009	16351
2008	14031
2007	13750
2003	7928
2006	7592
2005	5376
2002	5135
2000	4834
2010	4087

dtype: int64

- Looking at the interactions by year, the most are from the group that had a '0' entered in this column. While I didn't remove the 'year' column from this notebook altogether at the beginning or remove all songs that have the number '0' in the year column, this could be an interesting code to run in terms of how the year of the song could affect the recommendation system. The rest of the top 10 are years that are all in the decade of the 2000's except for the year 2010, so the newer the song the more likely an interaction will transpire.

- Removing the year column:

```
df.drop(['year'], axis=1, inplace=True)
```

- Removing all songs with a '0' in the year column: df_final=df_final[df_final.year>=1]

Interacted songs by artist

```
In [45]: most_rated = df_final.groupby('artist_name').size().sort_values(ascending = False)[:10]
most_rated
```

```
Out[45]:artist_name
```

Coldplay	5317
The Killers	4128
Florence + The Machine	2896
Kings Of Leon	2864
the bird and the bee	2387
LCD Soundsystem	2168
Vampire Weekend	2145
Justin Bieber	2130
Octopus Project	1825
Soltero	1691

dtype: int64

Interacted songs by release

```
In [46]: most_rated = df_final.groupby('release').size().sort_values(ascending = False)[:10]
most_rated
```

```
Out[46]:release
```

Ray Guns Are Not Just The Future	1843
My Worlds	1724
Vampire Weekend	1699
Hell Train	1691
Give Up	1604
Lungs	1511
A Rush Of Blood To The Head	1277
Only By The Night	1198
Save Me_San Francisco	1065
Identification Parade	1042

dtype: int64

- Was curious to see how the interactions were for the song, artist name, and album releases, and there is definitely a select few that are very popular with over 1000 interactions. I believe that could be a solid sign that the recommendation system is doing a good job in finding the right songs to recommend and that there is somewhat of a balance there between popular and more unknown songs.

Interacted songs by play counts

```
In [47]: most_rated = df_final.groupby('play_count').size().sort_values(ascending = False)[:10]
most_rated
```

Out[47]: play_count

1	72473
2	23890
3	10774
4	5874
5	4865

dtype: int64

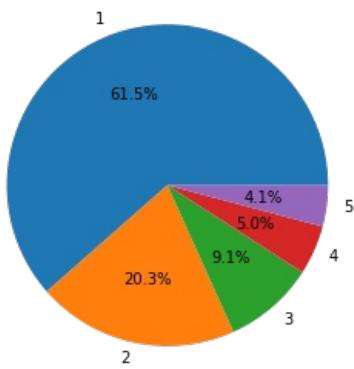
```
In [48]: df_final.play_count.value_counts()
```

Out[48]: 1 72473

2	23890
3	10774
4	5874
5	4865

Name: play_count, dtype: int64

```
In [49]: plt.pie(df_final.play_count.value_counts(), autopct = '%.1f%%', radius = 1.2, labels = ['1','2','3','4','5'])
plt.show()
```



- Showing the total number of interactions per play count which has been set to 5 maximum, the amount of interactions decreases as each number of play counts increases. Just one interaction has the highest number and 5 total interactions has the lowest.
- In the pie chart with the amount in percentages, 61.5 percent of the interactions occur with just one play count, followed by 20.3 percent with two play counts and only 4.1 percent with 5 play counts. So you're more likely than not going to have just one interaction between a particular user and song.

Observations and Insights:

- Despite a reduced dataset from the revised requirements put in place, there are still plenty of songs that have been played by users or have been seeked out by users to listen to and more than once. Even looking at the top 50 songs there is over 400 interactions per song.
- The total amount of interactions in the list of users is smaller than the song interactions with the top user haivng 247 interactions. With a million songs in the database this means that a stronger recommendation could be implemented so that more interactions take place.

Songs played in a year

```
In [50]: count_songs = df_final.groupby('year').count()['title']
```

```
count = pd.DataFrame(count_songs)
count.drop(count.index[0], inplace = True)
count.tail()
```

Out[50]: title

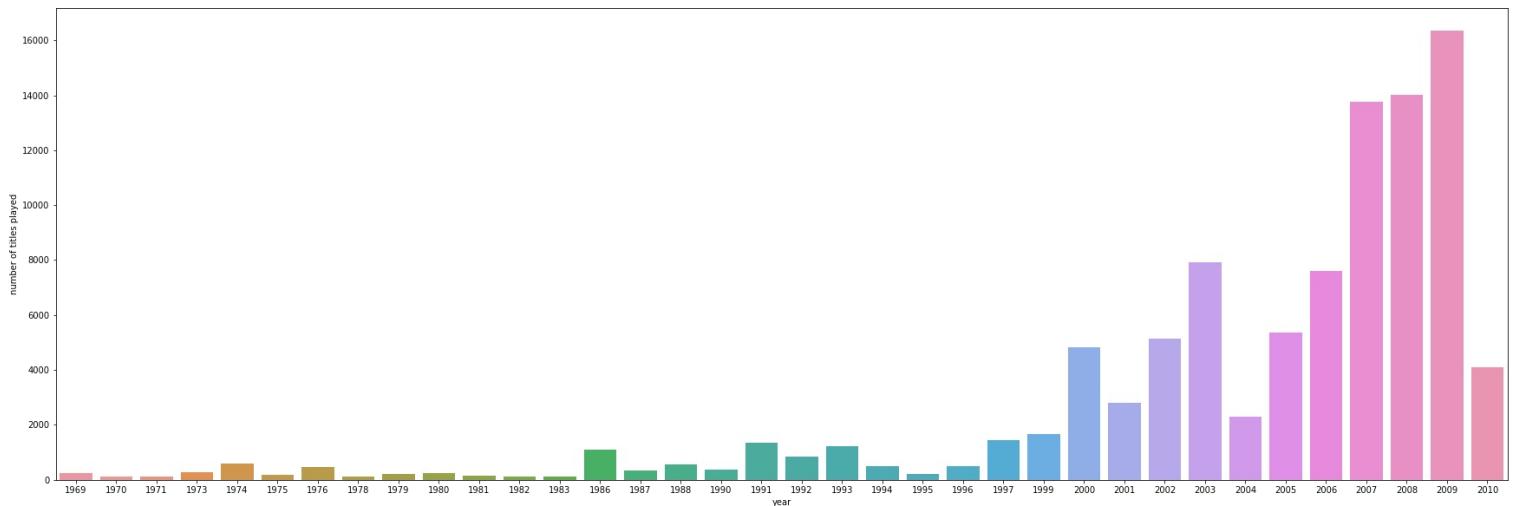
year	title
2006	7592
2007	13750
2008	14031
2009	16351
2010	4087

```
In [51]: plt.figure(figsize = (30, 10))
```

```
sns.barplot(x = count.index,
             y = 'title',
             data = count,
             estimator = np.median)

plt.ylabel('number of titles played')

plt.show()
```



Songs played by artist

```
In [52]: count_artist_names = df_final.groupby('artist_name').count()['title']

count = pd.DataFrame(count_artist_names)

count.drop(count.index[0], inplace = True)

count.head(10)
```

Out[52]:

title

artist_name	title
A Day To Remember	213
Adam Lambert	540
Aerosmith	403
Alicia Keys	362
Alliance Ethnik	717
Aloe Blacc	117
Amy Winehouse	1268
Angels and Airwaves	301
Arcade Fire	409
Avril Lavigne	156

Songs played by release

```
In [53]: count_artist_names = df_final.groupby('release').count()['title']

count = pd.DataFrame(count_artist_names)

count.drop(count.index[0], inplace = True)

count.tail(10)
```

Out[53]:

title

release

Workout Holiday	108
X & Y	503
Yellow	583
You Know I'm No Good	139
You Know I'm No Good - GFK Version	285
Youth And Young Manhood	132
Yuugata Generation	183
the bird and the bee	405
xx	120
À La Carte	119

- Observing the samples for interactions from releases and artists, it is a bit more lopsided from the artist sample than the release sample. I do

wonder if multiple playcount numbers from particular songs increase certain artists interaction totals and that is where the major difference is between artists.

Observations and Insights:

- From the last five years there is a noticeable difference in songs that came from those years compared to previous years. This might identify a certain demographic in who is using the database and what songs and artists they are listening to.
- Looking at the graph it clearly shows that the amount of songs in the database decrease the farther back you go in years. From this century the majority of the songs come from that time frame, while there is some songs from the 1990's but much lower in quantity. It becomes very small in numbers before 1990, so the newer the song the more likely that it is going to be listened to by the user.

Think About It: What other insights can be drawn using exploratory data analysis?

- From the values you can see what users are listening to more than anything which is newer music and that they are not seeking out older music, but it could also depend on the age of the user.
- If the majority of songs are newer, then building a recommendation system based off that information could make it easier to find the 10 best recommended songs depending on past interactions.
- If there are less unique artists then this could also be important to constructing an algorithm where it can locate popular artists but also find lesser known artists that can be recommended to users.

```
In [54]: from google.colab import files  
df_final.to_csv('df_final.csv', encoding = 'utf-8-sig')  
files.download('df_final.csv')
```

Milestone 2

- Now that we have explored the data, let's apply different algorithms to build recommendation systems.
- Note: Use the shorter version of the data, i.e., the data after the cutoffs as used in Milestone 1.

Load the dataset

```
In [55]: from google.colab import drive  
drive.mount('/content/drive')
```

Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount("/content/drive", force_remount=True).

```
In [56]: import warnings  
warnings.filterwarnings('ignore')
```

```
import numpy as np  
import pandas as pd  
  
import matplotlib.pyplot as plt  
  
import seaborn as sns  
  
from sklearn.metrics.pairwise import cosine_similarity  
  
from collections import defaultdict  
  
from sklearn.metrics import mean_squared_error  
  
!pip install scipy  
  
from pandas.core.arrays.interval import le  
from sklearn.preprocessing import LabelEncoder  
from sklearn.metrics.pairwise import cosine_similarity
```

```
Looking in indexes: https://pypi.org/simple, https://us-python.pkg.dev/colab-wheels/public/simple/  
Requirement already satisfied: scipy in /usr/local/lib/python3.8/dist-packages (1.7.3)  
Requirement already satisfied: numpy<1.23.0,>=1.16.5 in /usr/local/lib/python3.8/dist-packages (from scipy) (1.21.6)  
In [57]: df_final = pd.read_csv('/content/drive/MyDrive/df_final (6).csv')
```

Popularity-Based Recommendation Systems

Let's take the count and sum of play counts of the songs and build the popularity recommendation systems based on the sum of play counts.

```
In [58]: average_count = df_final.groupby(['song_id']).mean()['play_count']
```

```
play_freq = df_final.groupby(['song_id']).count()['play_count']
```

```
In [59]: final_play = pd.DataFrame({'avg_count': average_count, 'play_freq': play_freq})  
  
final_play.head()
```

```
Out[59]: avg_count play_freq
```

song_id	avg_count	play_freq
21	1.622642	265
22	1.492424	132
52	1.729216	421
62	1.728070	114
93	1.452174	115

Now, let's create a function to find the top n songs for a recommendation based on the average play count of song. We can also add a threshold for a minimum number of playcounts for a song to be considered for recommendation.

In [60]: `def top_n_songs(data, n, min_interaction = 100):`

```
    recommendations = data[data['play_freq'] > min_interaction]
    recommendations = recommendations.sort_values(by = 'avg_count', ascending = False)
    return recommendations.index[:n]
```

In [61]: `list(top_n_songs(final_play, 10, 100))`

Out[61]: [7224, 6450, 9942, 5531, 5653, 8483, 2220, 657, 614, 352]

In [62]: `df_final[df_final['song_id'] == 7224]`

Out[62]:

	Unnamed: 0	user_id	song_id	play_count	title	release	artist_name	year
1850	24098	23297	7224	5	Victoria (LP Version)	Hit By A Train: The Best Of Old 97's	Old 97's	2006
2037	28161	76307	7224	1	Victoria (LP Version)	Hit By A Train: The Best Of Old 97's	Old 97's	2006
2245	31860	26149	7224	1	Victoria (LP Version)	Hit By A Train: The Best Of Old 97's	Old 97's	2006
4362	76415	33280	7224	5	Victoria (LP Version)	Hit By A Train: The Best Of Old 97's	Old 97's	2006
4766	81358	42453	7224	1	Victoria (LP Version)	Hit By A Train: The Best Of Old 97's	Old 97's	2006
...
111323	1883439	31269	7224	5	Victoria (LP Version)	Hit By A Train: The Best Of Old 97's	Old 97's	2006
112158	1903542	28099	7224	5	Victoria (LP Version)	Hit By A Train: The Best Of Old 97's	Old 97's	2006
114439	1948494	36058	7224	1	Victoria (LP Version)	Hit By A Train: The Best Of Old 97's	Old 97's	2006
115786	1968497	27606	7224	5	Victoria (LP Version)	Hit By A Train: The Best Of Old 97's	Old 97's	2006
115887	1969785	902	7224	5	Victoria (LP Version)	Hit By A Train: The Best Of Old 97's	Old 97's	2006

107 rows × 8 columns

- Using the top song (7224) based on average play count, this code shows some of its users who interacted with the song.

User User Similarity-Based Collaborative Filtering

To build the user-user-similarity-based and subsequent models we will use the "surprise" library.

In [63]: `!pip install surprise`

```
Looking in indexes: https://pypi.org/simple, https://us-python.pkg.dev/colab-wheels/public/simple/
Requirement already satisfied: surprise in /usr/local/lib/python3.8/dist-packages (0.1)
Requirement already satisfied: scikit-surprise in /usr/local/lib/python3.8/dist-packages (from surprise) (1.1.3)
Requirement already satisfied: joblib>=1.0.0 in /usr/local/lib/python3.8/dist-packages (from scikit-surprise->surprise) (1.2.0)
Requirement already satisfied: numpy>=1.17.3 in /usr/local/lib/python3.8/dist-packages (from scikit-surprise->surprise) (1.21.6)
Requirement already satisfied: scipy>=1.3.2 in /usr/local/lib/python3.8/dist-packages (from scikit-surprise->surprise) (1.7.3)
```

In [64]: `from surprise import accuracy`

```
from surprise.reader import Reader
from surprise.dataset import Dataset
from surprise.model_selection import GridSearchCV
from surprise.model_selection import train_test_split
from surprise.prediction_algorithms.knns import KNNBasic
from surprise.prediction_algorithms.matrix_factorization import SVD
from surprise.model_selection import KFold
from surprise import CoClustering
```

Some useful functions

Below is the function to calculate precision@k and recall@k, RMSE and F1_Score@k to evaluate the model performance.

Think About It: Which metric should be used for this problem to compare different models?

- When building a recommendation system that will be implementing different models I do think the F1 score should be used because it will show the mean between the precision and recall. While I do believe that the precision and recall metrics are important in the beginning stages, the F1 score will give a better idea of how its performing while comparing models during this assessment.

In [65]: `def precision_recall_at_k(model, k = 30, threshold = 1.5):`

```

"""Return precision and recall at k metrics for each user"""

user_est_true = defaultdict(list)

predictions=model.test(testset)

for uid, _, true_r, est, _ in predictions:
    user_est_true[uid].append((est, true_r))

precisions = dict()
recalls = dict()
for uid, user_ratings in user_est_true.items():

    user_ratings.sort(key = lambda x : x[0], reverse = True)

    n_rel = sum((true_r >= threshold) for _, true_r in user_ratings)

    n_rec_k = sum((est >= threshold) for (est, _) in user_ratings[ : k])

    n_rel_and_rec_k = sum(((true_r >= threshold) and (est >= threshold))
                           for (est, true_r) in user_ratings[ : k])

    precisions[uid] = n_rel_and_rec_k / n_rec_k if n_rec_k != 0 else 0

    recalls[uid] = n_rel_and_rec_k / n_rel if n_rel != 0 else 0

precision = round((sum(prec for prec in precisions.values()) / len(precisions)), 3)

recall = round((sum(rec for rec in recalls.values()) / len(recalls)), 3)

accuracy.rmse(predictions)

print('Precision: ', precision)

print('Recall: ', recall)

print('F_1 score: ', round((2 * precision * recall) / (precision + recall), 3))

```

Think About It: In the function precision_recall_at_k above the threshold value used is 1.5. How precision and recall are affected by changing the threshold? What is the intuition behind using the threshold value of 1.5?

- I think it is fair to have the threshold at 1.5 because it allows the model to be reduced so it can locate the 10 best recommendations for the user. Because it tightens things up a bit I do think it could impact the precision and recall but possibly for the better. If it is lower than around 0.6 or so then it could indicate that the recommendation system is not working properly and the threshold number should be changed. Setting the threshold value of 1.5 is a good spot in that it rests somewhere within the average count values and also looking at the past history regarding the play counts of songs.

In [66]: reader = Reader(rating_scale= (0, 5))

```

data = Dataset.load_from_df(df_final[['user_id', 'song_id', 'play_count']], reader)

trainset, testset = train_test_split(data, test_size=0.4, random_state = 42)

```

Think About It: How changing the test size would change the results and outputs?

- I think it would be because a larger portion of your data should be used toward the model and using it for fitting and training. So if there is a higher test size then it is taking away a larger amount of the data that needs to be used for the model training. This in turn could then affect the outputs and results of the data.

In [67]: sim_options = {'name': 'msd',
 'user_based': True}

sim_user_user = KNNBasic(sim_options = sim_options, k = 30, random_state = 1, verbose = False)

sim_user_user.fit(trainset)

precision_recall_at_k(sim_user_user)

RMSE: 1.0672

Precision: 0.412

Recall: 0.598

F_1 score: 0.488

Observations and Insights:

- The precision value looks a little low which at this juncture may indicate that the song recommendations provided are not very relevant.
- The recall value is right at 0.6 which means that relevant songs are being recommended to the user and that the model has the value at a good spot.
- The RMSE is still over 1 and that could be improved so the predicted ratings are closer to the actual ratings. The overall F1 score is still lower than it should be so this could also be improved with a better model from the data.

In [68]: sim_user_user.predict(6958, 1671, r_ui = 2, verbose = True)

```

user: 6958      item: 1671      r_ui = 2.00      est = 1.63      {'actual_k': 30, 'was_impossible': False}
Out[68]: Prediction(uid=6958, iid=1671, r_ui=2, est=1.6254190211665536, details={'actual_k': 30, 'was_impossible': False})

```

In [69]:

```
sim_user_user.predict(6958, 3232, verbose = True)

user: 6958      item: 3232      r_ui = None    est = 1.36  {'actual_k': 30, 'was_impossible': False}
Out[69]:Prediction(uid=6958, iid=3232, r_ui=None, est=1.3609600600037504, details={'actual_k': 30, 'was_impossible': False})
```

Observations and Insights:

- Comparing the two samples it looks like the average play count that is predicted for the user with a listened song is considerably higher than the same user that hasn't listened to the particular song. With the actual rating being 2 for the user that has listened to the song, the predicted rating is not too far off but could be improved.

Now, let's try to tune the model and see if we can improve the model performance.

```
In [70]: param_grid = {'k': [10, 20, 30], 'min_k': [3, 6, 9],
                     'sim_options': {'name': ['cosine', 'pearson', "pearson_baseline"],
                                    'user_based': [True], "min_support": [2, 4]}}
}
```

```
gs = GridSearchCV(KNNBasic, param_grid, measures = ['rmse'], cv = 3, n_jobs = -1)
```

```
gs.fit(data)
```

```
print(gs.best_score['rmse'])
```

```
1.0473120682571855
```

```
In [71]: print(gs.best_params['rmse'])
```

```
{'k': 30, 'min_k': 9, 'sim_options': {'name': 'pearson_baseline', 'user_based': True, 'min_support': 2}}
```

```
In [72]: sim_options = {'name': 'pearson_baseline',
                      'user_based': True}
```

```
sim_user_user_optimized = KNNBasic(sim_options = sim_options, k = 30, min_k = 9, random_state = 1, verbose = False)
```

```
sim_user_user_optimized.fit(trainset)
```

```
precision_recall_at_k(sim_user_user_optimized)
```

```
RMSE: 1.0521
```

```
Precision: 0.413
```

```
Recall: 0.721
```

```
F_1 score: 0.525
```

Observations and Insights:

- The noticeable change from the metrics was the recall that improved after the hyperparameters were fine tuned. While they were all ready at a good level before the model was tuned, it showed the most improvement.
- The RMSE metric was lowered but not by very much, so it could be a better fit but the fine tuning helped. It should be improved though so the accuracy of the model is better for predictions. The F1 score also improved slightly but mainly from the improvement in the recall value. This is another aspect that could be increased closer to 1 to enhance the accuracy of the model.
- The one metric that stayed pretty much the same after fine tuning was the precision metric which barely improved, so because of this the F1 score did not change very much.

```
In [73]: sim_user_user_optimized.predict(6958, 1671, r_ui = 2, verbose = True)
```

```
user: 6958      item: 1671      r_ui = 2.00    est = 1.96  {'actual_k': 24, 'was_impossible': False}
Out[73]:Prediction(uid=6958, iid=1671, r_ui=2, est=1.962926073914969, details={'actual_k': 24, 'was_impossible': False})
```

```
In [74]: sim_user_user_optimized.predict(6958, 3232, verbose = True)
```

```
user: 6958      item: 3232      r_ui = None    est = 1.45  {'actual_k': 10, 'was_impossible': False}
Out[74]:Prediction(uid=6958, iid=3232, r_ui=None, est=1.4516261428486725, details={'actual_k': 10, 'was_impossible': False})
```

Observations and Insights:

- After optimizing and fine tuning the predictions the predicted rating for the user that has listened to a particular song (item 1671) increased substantially and is very close to the actual rating of 2. For the user that hasn't listened to a particular song (item 3232), the predicted values also increased from before the fine tuning but not by much.
- The k quantities were lower than 30 for both predictions after the optimizing, possibly meaning that there are less observations that are similar to them.

Think About It: Along with making predictions on listened and unknown songs can we get 5 nearest neighbors (most similar) to a certain song?

- I think using the KNN algorithm could potentially be a solution to finding the 5 closest neighbors to a particular song. This way could calculate the distances among those points to determine what is the nearest 5 neighbors to that song in terms of similarity.

```
In [75]: sim_user_user_optimized.get_neighbors(0, k = 5)
```

```
Out[75]:[42, 1131, 17, 186, 249]
```

Below we will be implementing a function where the input parameters are:

data: A song dataset

user_id: A user-id against which we want the recommendations

top_n: The number of songs we want to recommend

algo: The algorithm we want to use for predicting the play_count

The output of the function is a set of top_n items recommended for the given user_id based on the given algorithm

In [76]: **def** get_recommendations(data, user_id, top_n, algo):

```
    recommendations = []

    user_item_interactions_matrix = data.pivot_table(index = 'user_id', columns = 'song_id', values = 'play_count')

    non_interacted_products = user_item_interactions_matrix.loc[user_id][user_item_interactions_matrix.loc[user_id] == 0]

    for item_id in non_interacted_products:

        est = algo.predict(user_id, item_id).est

        recommendations.append((item_id, est))

    recommendations.sort(key = lambda x : x[1], reverse = True)

    return recommendations[:top_n]
```

In [77]: recommendations = get_recommendations(df_final, 6958, 5, sim_user_user_optimized)

In [78]: pd.DataFrame(recommendations, columns = ['song_id', 'predicted_ratings'])

Out[78]:

	song_id	predicted_ratings
0	5531	2.553335
1	317	2.518269
2	4954	2.406776
3	8635	2.396606
4	5943	2.390723

Observations and Insights:

- The top 5 recommended songs from the new function all have predicted ratings that are well above the threshold of 2 that was used in previous algorithms and models. I believe this is a good sign that this new function is strong with its new predictions and have provided highly rated song recommendations to the user.

Correcting the play_counts and Ranking the above songs

In [79]: **def** ranking_songs(recommendations, final_rating):
 ranked_songs = final_rating.loc[[items[0] **for** items **in** recommendations]].sort_values('play_freq', ascending = False)
 ranked_songs = ranked_songs.merge(pd.DataFrame(recommendations, columns = ['song_id', 'predicted_ratings']), on = 'song_id')
 ranked_songs['corrected_ratings'] = ranked_songs['predicted_ratings'] - 1 / np.sqrt(ranked_songs['play_freq'])
 ranked_songs = ranked_songs.sort_values('corrected_ratings', ascending = False)

return ranked_songs

Think About It: In the above function to correct the predicted play_count a quantity $1/\sqrt{n}$ is subtracted. What is the intuition behind it? Is it also possible to add this quantity instead of subtracting?

- I think the reason behind the quantity being subtracted is that it allows the predicted play counts to have a more realistic and better prediction to the songs. Adding to the quantity might not condense the amount of similar items so that it finds the best recommendations, so perhaps that is the reason that it is better to subtract the quantity.

In [80]: ranking_songs(recommendations, final_play)

Out[80]:

	song_id	play_freq	predicted_ratings	corrected_ratings
0	5531	618	2.553335	2.513109
2	317	411	2.518269	2.468943
1	5943	423	2.390723	2.342101
3	4954	183	2.406776	2.332854
4	8635	155	2.396606	2.316284

Observations and Insights:

- The rankings of the predicted ratings and corrected ratings for the top 5 are both parallel in that they are ranked in order for each column. What also is relevant here is that the play frequency is also parallel with the ratings in that they are ranked in order like the other columns. This indicates that there is a correlation between how much a song is played and how high the ratings will be for that song. The impact of the user-item interaction then can't be understated and that the play frequency is important to the ratings and makes it more dependable in its predictions.

Item Item Similarity-based collaborative filtering recommendation systems

In [81]: sim_options = {'name': 'cosine',

```
'user_based': False}

sim_item_item = KNNBasic(sim_options = sim_options, random_state = 1, verbose = False)

sim_item_item.fit(trainset)

precision_recall_at_k(sim_item_item)
```

RMSE: 1.0394
 Precision: 0.307
 Recall: 0.562
 F_1 score: 0.397

Observations and Insights:

- After running the model it looks like the precision metric is quite low which is keeping the F1 score much lower than it should be. This must mean that the recommendations provided for the items are not relevant at this point. The recall is a bit lower than it should be which is around 0.6 so improvement is needed here as well where the relevant items recommended could be fine tuned a bit.
- The RMSE metric is over 1 which means that this is another area that can be improved since it is still far from being accurate.

In [82]: sim_item_item.predict(6958, 1671, r_ui = 2, verbose = True)

```
user: 6958      item: 1671      r_ui = 2.00  est = 1.36  {'actual_k': 20, 'was_impossible': False}
Out[82]: Prediction(uid=6958, iid=1671, r_ui=2, est=1.3614157231762556, details={'actual_k': 20, 'was_impossible': False})
```

In [83]: sim_item_item.predict(3232, 1671, verbose = True)

```
user: 3232      item: 1671      r_ui = None   est = 1.70  {'was_impossible': True, 'reason': 'User and/or item is unknown.'}
Out[83]: Prediction(uid=3232, iid=1671, r_ui=None, est=1.6989607635206787, details={'was_impossible': True, 'reason': 'User and/or item is unknown.'})
```

Observations and Insights:

- The estimated prediction of 1.36 for the user that has listened to this particular song is much lower than the provided threshold of 2 with the k value being around 20. Maybe this means that with less likely neighbors that this is not a good song to recommend based off the prediction.
- The estimated prediction of 1.70 with the user that has not listened to the particular song is much higher than the listened to song. It is still below the threshold of 2 which means at this point the metrics for recommendations are still not very strong.

In [84]: param_grid = {'k': [10, 20, 30], 'min_k': [3, 6, 9],
 'sim_options': {'name': ["cosine", 'pearson', "pearson_baseline"]},
 'user_based': [False], "min_support": [2, 4]}

```
gs = GridSearchCV(KNNBasic, param_grid, measures = ['rmse'], cv = 3, n_jobs = -1)
gs.fit(data)

print(gs.best_score['rmse'])
```

1.024025696240705

In [85]: print(gs.best_params['rmse'])

```
{'k': 30, 'min_k': 6, 'sim_options': {'name': 'pearson_baseline', 'user_based': False, 'min_support': 2}}
```

Think About It: How do the parameters affect the performance of the model? Can we improve the performance of the model further? Check the list of hyperparameters here.

- The usage of the grid search is vital because it will take the hyperparameters listed in the above model and look at all possibilities within those hyperparameters and models to construct a model that will be most beneficial. Perhaps adjusting the values in the parameter grid could be one way of making the model better depending on how the final output looks after entering the data.

In [86]: sim_options = {'name': 'msd',
 'user_based': False}

```
sim_item_item_optimized = KNNBasic(sim_options = sim_options, k = 30, min_k = 6, random_state = 1, verbose = False)
sim_item_item_optimized.fit(trainset)

precision_recall_at_k(sim_item_item_optimized)
```

RMSE: 1.0423
 Precision: 0.34
 Recall: 0.563
 F_1 score: 0.424

Observations and Insights:

- After optimizing the model it looks as though that the tuned hyperparameters did not change the metrics too much. The RMSE metric even regressed and is slightly farther away from being accurate. The precision metric improved more than the other metrics but it was not a significant increase, so it helped improve the F1 score but only slightly because the recall metric stayed roughly the same.
- Overall the accuracy of the metrics is still well below par so the fine tuning of the model did not have much of an effect on the metrics. Maybe going back and revising the parameters in the model could help these values improve.

In [87]: sim_item_item_optimized.predict(6958, 1671, r_ui = 2, verbose = True)

```
user: 6958      item: 1671      r_ui = 2.00  est = 1.33  {'actual_k': 20, 'was_impossible': False}
Out[87]: Prediction(uid=6958, iid=1671, r_ui=2, est=1.3319659681769878, details={'actual_k': 20, 'was_impossible': False})
```

In [88]: sim_item_item_optimized.predict(6958, 3232, verbose = True)

```
user: 6958      item: 3232      r_ui = None    est = 1.47    {'actual_k': 20, 'was_impossible': False}
Out[88]:Prediction(uid=6958, iid=3232, r_ui=None, est=1.4746440368585654, details={'actual_k': 20, 'was_impossible': False})
```

Observations and Insights:

- After optimizing the model it looks like both predictions for the user went down and were farther away from the threshold of 2. This is on par with the other outputs after fine tuning the hyperparameters in that they had little impact and even showed signs of regression in some areas. The song that was listened to by the user stayed roughly the same with little regression while the song that was not listened to showed a lot more decline in its estimated rating. Both outputs had actual k values of 20 like the previous model before the optimization.

```
In [89]: sim_item_item_optimized.get_neighbors(0, k = 5)
```

```
Out[89]:[3, 10, 24, 30, 36]
```

```
In [90]: recommendations = get_recommendations(df_final, 6958, 5, sim_item_item_optimized)
```

```
In [91]: pd.DataFrame(recommendations, columns = ['song_id', 'predicted_play_count'])
```

```
Out[91]:   song_id  predicted_play_count
```

0	9942	2.048219
1	2842	1.949116
2	3050	1.885211
3	4939	1.771190
4	1691	1.743513

```
In [92]: ranking_songs(recommendations, final_play)
```

```
Out[92]:   song_id  play_freq  predicted_ratings  corrected_ratings
```

3	9942	150	2.048219	1.966570
2	2842	232	1.949116	1.883463
1	3050	233	1.885211	1.819698
4	4939	133	1.771190	1.684479
0	1691	249	1.743513	1.680140

Observations and Insights:

- From the 5 selected neighbors provided the model recommended 5 songs that are all above or at least somewhat close to the threshold of 2. The play frequencies for the songs are all high enough to show that this is a fairly reliable prediction for the 5 songs. The corrected ratings are all pretty close to the predicted ratings, so this shows that this is also a good model for recommending songs if you are looking simply at the closest neighbors to the user.

Model Based Collaborative Filtering - Matrix Factorization

Model-based Collaborative Filtering is a personalized recommendation system, the recommendations are based on the past behavior of the user and it is not dependent on any additional information. We use latent features to find recommendations for each user.

```
In [93]: svd = SVD(random_state = 1)
```

```
    svd.fit(trainset)
    precision_recall_at_k(svd)
```

```
RMSE: 1.0252
```

```
Precision: 0.41
```

```
Recall: 0.633
```

```
F_1 score: 0.498
```

```
In [94]: svd.predict(6958, 1671, r_ui = 2, verbose = True)
```

```
user: 6958      item: 1671      r_ui = 2.00    est = 1.27    {'was_impossible': False}
```

```
Out[94]:Prediction(uid=6958, iid=1671, r_ui=2, est=1.267473397214638, details={'was_impossible': False})
```

```
In [95]: svd.predict(6958, 3232, verbose = True)
```

```
user: 6958      item: 3232      r_ui = None    est = 1.56    {'was_impossible': False}
```

```
Out[95]:Prediction(uid=6958, iid=3232, r_ui=None, est=1.5561675084403663, details={'was_impossible': False})
```

Improving matrix factorization based recommendation system by tuning its hyperparameters

```
In [96]: param_grid = {'n_epochs': [10, 20, 30], 'lr_all': [0.001, 0.005, 0.01],
    'reg_all': [0.2, 0.4, 0.6]}
```

```
    gs = GridSearchCV(SVD, param_grid, measures = ['rmse'], cv = 3, n_jobs = -1)
```

```
    gs.fit(data)
```

```
    print(gs.best_score['rmse'])
```

```
1.0121087151082329
```

```
In [97]: print(gs.best_params['rmse'])
```

```
{'n_epochs': 30, 'lr_all': 0.01, 'reg_all': 0.2}
```

Think About It: How do the parameters affect the performance of the model? Can we improve the performance of the model further? Check the available hyperparameters here.

- The parameters that they are all set at it looks like it has improved the metrics considerably in most areas. The F1 score has improved but is still below where it should be, while the RMSE metric has been lowered but only slightly. The prediction value for the user that has listened to a particular song is the only noticeable regression. None of the parameters that are being used are near the default options, so perhaps using those values might be a better method to see how the model has affected the outputs.

```
In [98]: svd_optimized = SVD(n_epochs = 30, lr_all = 0.01, reg_all = 0.2, random_state = 1)

    svd_optimized = svd_optimized.fit(trainset)

    precision_recall_at_k(svd_optimized)
```

RMSE: 1.0141
Precision: 0.415
Recall: 0.635
F_1 score: 0.502

Observations and Insights:

- Using the best possible parameters from the grid search it has shown to have made the metrics better in all four areas. Maybe using the default options would have not made the model better and that this is the best model for future recommendations after through all the possible iterations. It has pushed the F1 score closer to 0.6 which will make this a more reliable model.

```
In [99]: svd_optimized.predict(6958, 1671, r_ui = 2, verbose = True)
```

user: 6958 item: 1671 r_ui = 2.00 est = 1.34 {'was_impossible': False}
Out[99]: Prediction(uid=6958, iid=1671, r_ui=2, est=1.3432395286125096, details={'was_impossible': False})

```
In [100]: svd_optimized.predict(6958, 3232, verbose = True)
```

user: 6958 item: 3232 r_ui = None est = 1.44 {'was_impossible': False}
Out[100]: Prediction(uid=6958, iid=3232, r_ui=None, est=1.442548446117648, details={'was_impossible': False})

Observations and Insights:

- The optimized hyperparameters increased the estimated value by a small amount with using 2 still as a threshold with a particular song that the user has listened to. The song that the user has not listened to has an estimated value that decreased after fine tuning the hyperparameters. I do wonder if this process that uses latent features is affecting the values after tuning the hyperparameters or any regression is more of an aberration.

```
In [101]: svd_recommendations = get_recommendations(df_final, 6958, 5, svd_optimized)
```

```
In [102]: ranking_songs(svd_recommendations, final_play)
```

```
Out[102]:   song_id  play_freq  predicted_ratings  corrected_ratings

      2      7224        107       2.601899       2.505225
      1      5653        108       2.108728       2.012502
      4      8324        96       2.014091       1.912029
      0      9942        150       1.940115       1.858465
      3      6450        102       1.952493       1.853478
```

Observations and Insights:

- The tuned hyperparameters have I think changed the model for the better, with the top 5 song recommendations having higher ratings in both columns. They are all a lot closer to the 2 threshold with all the songs ratings being above or just below that threshold value. The difference between the predicted and corrected ratings look about the same after the model was tuned. All the songs have high enough play frequencies that this recommendation model looks stronger.

Cluster Based Recommendation System

In clustering-based recommendation systems, we explore the similarities and differences in people's tastes in songs based on how they rate different songs. We cluster similar users together and recommend songs to a user based on play_counts from other users in the same cluster.

```
In [103]: clust_baseline = CoClustering(random_state = 1)
```

```
    clust_baseline.fit(trainset)

    precision_recall_at_k(clust_baseline)
```

RMSE: 1.0487
Precision: 0.397
Recall: 0.582
F_1 score: 0.472

```
In [104]: clust_baseline.predict(6958, 1671, r_ui = 2, verbose = True)
```

user: 6958 item: 1671 r_ui = 2.00 est = 1.29 {'was_impossible': False}
Out[104]: Prediction(uid=6958, iid=1671, r_ui=2, est=1.2941824757363074, details={'was_impossible': False})

```
In [105]: clust_baseline.predict(6958, 3232, verbose = True)
```

user: 6958 item: 3232 r_ui = None est = 1.48 {'was_impossible': False}
Out[105]: Prediction(uid=6958, iid=3232, r_ui=None, est=1.4785259100797417, details={'was_impossible': False})

Improving clustering-based recommendation system by tuning its hyper-parameters

```
In [106]: param_grid = {'n_cltr_u': [5, 6, 7, 8], 'n_cltr_i': [5, 6, 7, 8], 'n_epochs': [10, 20, 30]}

gs = GridSearchCV(CoClustering, param_grid, measures = ['rmse'], cv = 3, n_jobs = -1)

gs.fit(data)

print(gs.best_score['rmse'])

1.061515101753293
```

In [107]: print(gs.best_params['rmse'])

{'n_cltr_u': 5, 'n_cltr_i': 5, 'n_epochs': 10}

Think About It: How do the parameters affect the performance of the model? Can we improve the performance of the model further? Check the available hyperparameters here.

- I think that having less clusters for the items and users will benefit the recommendation system in that it will make it easier to recommend items and help find users that are similar. Having too many clusters I wonder could make things too complicated for the recommendation system and would require too many iterations of the optimization loop.

```
In [108]: clust_tuned = CoClustering(n_cltr_u = 3, n_cltr_i = 3, n_epochs = 30, random_state = 1)
```

```
clust_tuned.fit(trainset)

precision_recall_at_k(clust_tuned)
```

RMSE: 1.0487
Precision: 0.397
Recall: 0.582
F_1 score: 0.472

Observations and Insights:

- Condensing the total number of clusters to the default values of 3 after the tuning, the F1 score is below average with the recall being close to average and the precision being relatively low. The RMSE metric is around 1 which is a bit high but isn't too far off from other past models that have been implemented. I wonder if more clusters would help these metrics scores at all.

```
In [109]: clust_tuned.predict(6958, 1671, r_ui = 2, verbose = True)
```

```
user: 6958 item: 1671 r_ui = 2.00 est = 1.29 {'was_impossible': False}
Out[109]: Prediction(uid=6958, iid=1671, r_ui=2, est=1.2941824757363074, details={'was_impossible': False})
```

```
In [110]: clust_tuned.predict(6958, 3232, verbose = True)
```

```
user: 6958 item: 3232 r_ui = None est = 1.48 {'was_impossible': False}
Out[110]: Prediction(uid=6958, iid=3232, r_ui=None, est=1.4785259100797417, details={'was_impossible': False})
```

Observations and Insights:

- It looks as though the clusters have made the estimated values for both predictions with the user-song interaction a bit below average. The cluster with the user having not listened to the particular song is definitely higher than the song that the user has listened to using the threshold of 2. This seems to be not too much of a surprise with the metrics being below average overall as well. There is definitely some room for improvements regarding the fine tuning of the hyperparameters with the clustering.

Implementing the recommendation algorithm based on optimized CoClustering model

```
In [111]: def get_recommendations(data, user_id, top_n, algo):
```

```
    recommendations = []

    user_item_interactions_matrix = data.pivot(index = 'user_id', columns = 'song_id', values = 'play_count')

    non_interacted_products = user_item_interactions_matrix.loc[user_id][user_item_interactions_matrix.loc[user_id].isna()]

    for item_id in non_interacted_products:
        est = algo.predict(user_id, item_id).est
        recommendations.append((item_id, est))

    recommendations.sort(key = lambda x: x[1], reverse = True)

    return recommendations[:top_n]
```

```
In [112]: clustering_recommendations = get_recommendations(df_final, 6958, 5, clust_tuned)
```

Correcting the play_count and Ranking the above songs

```
In [113]: def ranking_songs(recommendations, final_rating):
```

```
    ranked_songs = final_play.loc[[items[0] for items in recommendations]].sort_values('play_freq', ascending = False)

    ranked_songs = ranked_songs.merge(pd.DataFrame(recommendations, columns = ['song_id', 'predicted_ratings']))

    ranked_songs['corrected_ratings'] = ranked_songs['predicted_ratings'] - 1 / np.sqrt(ranked_songs['play_freq'])

    ranked_songs = ranked_songs.sort_values('corrected_ratings', ascending = False)

    return ranked_songs
```

```
In [114]:
```

```
ranking_songs(clustering_recommendations, final_play)
```

```
Out[114]:   song_id  play_freq  predicted_ratings  corrected_ratings
```

2	7224	107	3.094797	2.998124
4	8324	96	2.311498	2.209436
1	9942	150	2.215039	2.133390
0	5531	618	2.124563	2.084337
3	4831	97	2.123783	2.022248

Observations and Insights:

- After going through the extraction and finding the songs that have not been interacted with, it looks like the ratings for these songs are even higher than past models. The clustering method then makes me wonder if this is a better way to help find songs for users in the recommendation system stemming from the results. The top 5 recommended songs also have high play frequencies, so I feel more comfortable using this model than the other models used as of right now.

Content Based Recommendation Systems

Think About It: So far we have only used the play_count of songs to find recommendations but we have other information/features on songs as well. Can we take those song features into account?

- I do think that other aspects of these songs can be used toward building a stronger recommendation system. If we know the genre of music or even the year of the album release then I think that can help identify user interests and that in turn can find similar users which can help recommend songs that the user would more likely listen to. The only part of this that could possibly hinder the best recommendation system regarding song information would be if any of those features are missing.

```
In [115]: df_small = df_final
```

```
In [116]: df_small['text'] = df_small['title'] + ' ' + df_small['release'] + ' ' + df_small['artist_name']
```

```
df_small.head()
```

```
Out[116]:   Unnamed: 0  user_id  song_id  play_count          title    release  artist_name  year           text
```

0	200	6958	447	1	Daisy And Prudence	Distillation	Erin McKeown	2000	Daisy And Prudence Distillation Erin McKeown
1	202	6958	512	1	The Ballad of Michael Valentine	Sawdust	The Killers	2004	The Ballad of Michael Valentine Sawdust The Ki...
2	203	6958	549	1	I Stand Corrected (Album)	Vampire Weekend	Vampire Weekend	2007	I Stand Corrected (Album) Vampire Weekend Vamp...
3	204	6958	703	1	They Might Follow You	Tiny Vipers	Tiny Vipers	2007	They Might Follow You Tiny Vipers Tiny Vipers
4	205	6958	719	1	Monkey Man	You Know I'm No Good	Amy Winehouse	2007	Monkey Man You Know I'm No Good Amy Winehouse

```
In [117]: df_small = df_small[['user_id', 'song_id', 'play_count', 'title', 'text']]
```

```
df_small = df_small.drop_duplicates(subset = ['title'])
```

```
df_small = df_small.set_index('title')
```

```
df_small.head(5)
```

```
Out[117]:      user_id  song_id  play_count           text
```

	user_id	song_id	play_count	text
				title
				Daisy And Prudence
				6958 447 1 Daisy And Prudence Distillation Erin McKeown
				The Ballad of Michael Valentine
				6958 512 1 The Ballad of Michael Valentine Sawdust The Ki...
				I Stand Corrected (Album)
				6958 549 1 I Stand Corrected (Album) Vampire Weekend Vamp...
				They Might Follow You
				6958 703 1 They Might Follow You Tiny Vipers Tiny Vipers
				Monkey Man
				6958 719 1 Monkey Man You Know I'm No Good Amy Winehouse

```
In [118]: df_small.shape
```

```
Out[118]: (561, 4)
```

```
In [119]: indices = pd.Series(df_small.index)
```

```
indices[ : 5]
```

```
Out[119]:0          Daisy And Prudence
1  The Ballad of Michael Valentine
2      I Stand Corrected (Album)
```

```
3             They Might Follow You
4                         Monkey Man
Name: title, dtype: object
In [120]: import nltk

nltk.download("punkt")
nltk.download("stopwords")
nltk.download("wordnet")

import re

from nltk import word_tokenize

from nltk.stem import WordNetLemmatizer

from nltk.corpus import stopwords

from sklearn.feature_extraction.text import CountVectorizer, TfidfVectorizer

import nltk

nltk.download('omw-1.4')
```

```
[nltk_data] Downloading package punkt to /root/nltk_data...
[nltk_data]   Package punkt is already up-to-date!
[nltk_data] Downloading package stopwords to /root/nltk_data...
[nltk_data]   Package stopwords is already up-to-date!
[nltk_data] Downloading package wordnet to /root/nltk_data...
[nltk_data]   Package wordnet is already up-to-date!
[nltk_data] Downloading package omw-1.4 to /root/nltk_data...
[nltk_data]   Package omw-1.4 is already up-to-date!
```

```
Out[120]: True
```

We will create a function to pre-process the text data:

```
In [121]: def tokenize(text):
```

```
    text = re.sub(r"[^a-zA-Z]", " ", text.lower())
    tokens = word_tokenize(text)

    words = [word for word in tokens if word not in stopwords.words("english")]

    text_lems = [WordNetLemmatizer().lemmatize(lem).strip() for lem in words]

    return text_lems
```

```
In [122]: tfidf = TfidfVectorizer(tokenizer = tokenize)
```

```
song_tfidf = tfidf.fit_transform(df_small['text'].values).toarray()
```

```
In [123]: pd.DataFrame(song_tfidf)
```

```
Out[123]:      0   1   2   3   4   5   6   7   8   9   ...  1427  1428  1429  1430  1431  1432  1433  1434  1435  1436
0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  ...  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0
1  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  ...  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0
2  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  ...  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0
3  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  ...  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0
4  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  ...  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0
...
556 0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  ...  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0
557 0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  ...  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0
558 0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  ...  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0
559 0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  ...  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0
560 0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  ...  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0
```

561 rows × 1437 columns

```
In [124]: similar_songs = cosine_similarity(song_tfidf, song_tfidf)
```

```
similar_songs
```

```
Out[124]: array([[1., 0., 0., ..., 0., 0., 0.],
 [0., 1., 0., ..., 0., 0., 0.],
 [0., 0., 1., ..., 0., 0., 0.],
 ...,
```

```
[0., 0., 0., ..., 1., 0., 0.],
[0., 0., 0., ..., 0., 1., 0.],
[0., 0., 0., ..., 0., 0., 1.]])
```

Finally, let's create a function to find most similar songs to recommend for a given song.

In [125]: `def recommendations(title, similar_songs):`

```
recommended_songs = []

idx = indices[indices == title].index[0]

score_series = pd.Series(similar_songs[idx]).sort_values(ascending = False)

top_10_indexes = list(score_series.iloc[1 : 11].index)
print(top_10_indexes)

for i in top_10_indexes:
    recommended_songs.append(list(df_small.index)[i])

return recommended_songs
```

Recommending 10 songs similar to Learn To Fly

In [126]: `recommendations('Learn To Fly', similar_songs)`

```
[509, 234, 423, 345, 394, 370, 371, 372, 373, 375]
```

Out[126]:

```
['Everlong',
 'The Pretender',
 'Nothing Better (Album)',
 'From Left To Right',
 'Lifespan Of A Fly',
 'Under The Gun',
 'I Need A Dollar',
 'Feel The Love',
 'All The Pretty Faces',
 'Bones']
```

In [127]: `df_final.query("release == 'There Is Nothing Left To Lose'")`

Out[127]:

		Unnamed: 0	user_id	song_id	play_count	title	release	artist_name	year	text
676	8275	75901	1188	2	Learn To Fly	There Is Nothing Left To Lose	Foo Fighters	1999	Learn To Fly There Is Nothing Left To Lose Foo...	
1369	19706	9097	1188	1	Learn To Fly	There Is Nothing Left To Lose	Foo Fighters	1999	Learn To Fly There Is Nothing Left To Lose Foo...	
1535	20895	19193	1188	1	Learn To Fly	There Is Nothing Left To Lose	Foo Fighters	1999	Learn To Fly There Is Nothing Left To Lose Foo...	
2171	31768	26149	1188	2	Learn To Fly	There Is Nothing Left To Lose	Foo Fighters	1999	Learn To Fly There Is Nothing Left To Lose Foo...	
3315	54773	35737	1188	1	Learn To Fly	There Is Nothing Left To Lose	Foo Fighters	1999	Learn To Fly There Is Nothing Left To Lose Foo...	
...	
115373	1964010	35114	1188	1	Learn To Fly	There Is Nothing Left To Lose	Foo Fighters	1999	Learn To Fly There Is Nothing Left To Lose Foo...	
115586	1966225	21484	1188	1	Learn To Fly	There Is Nothing Left To Lose	Foo Fighters	1999	Learn To Fly There Is Nothing Left To Lose Foo...	
116883	1988416	74529	1188	2	Learn To Fly	There Is Nothing Left To Lose	Foo Fighters	1999	Learn To Fly There Is Nothing Left To Lose Foo...	
117299	1991996	14416	1188	1	Learn To Fly	There Is Nothing Left To Lose	Foo Fighters	1999	Learn To Fly There Is Nothing Left To Lose Foo...	
117546	1995070	9570	1188	1	Learn To Fly	There Is Nothing Left To Lose	Foo Fighters	1999	Learn To Fly There Is Nothing Left To Lose Foo...	

222 rows × 9 columns

In [128]: `df_final.query("artist_name == 'Foo Fighters'")`

Out[128]:

		Unnamed: 0	user_id	song_id	play_count	title	release	artist_name	year	text
259	1272	49549	6525	2	The Pretender	Echoes_Silence_Patience & Grace	Foo Fighters	2007	The Pretender Echoes_Silence_Patience & Grac...	
517	4486	7320	6525	1	The Pretender	Echoes_Silence_Patience & Grace	Foo Fighters	2007	The Pretender Echoes_Silence_Patience & Grac...	
676	8275	75901	1188	2	Learn To Fly	There Is Nothing Left To Lose	Foo Fighters	1999	Learn To Fly There Is Nothing Left To Lose Foo...	

760	8780	57932	6525	2	The Pretender	Echoes_Silence_Patience & Grace	Foo Fighters	2007	The Pretender Echoes_Silence_Patience & Grac...
1369	19706	9097	1188	1	Learn To Fly	There Is Nothing Left To Lose	Foo Fighters	1999	Learn To Fly There Is Nothing Left To Lose Foo...
...
117299	1991996	14416	1188	1	Learn To Fly	There Is Nothing Left To Lose	Foo Fighters	1999	Learn To Fly There Is Nothing Left To Lose Foo...
117348	1992111	14416	6525	1	The Pretender	Echoes_Silence_Patience & Grace	Foo Fighters	2007	The Pretender Echoes_Silence_Patience & Grac...
117375	1992181	14416	9249	1	Everlong	The Colour And The Shape (Special Edition)	Foo Fighters	1992	Everlong The Colour And The Shape (Special Edi...
117546	1995070	9570	1188	1	Learn To Fly	There Is Nothing Left To Lose	Foo Fighters	1999	Learn To Fly There Is Nothing Left To Lose Foo...
117703	1995402	9570	9249	1	Everlong	The Colour And The Shape (Special Edition)	Foo Fighters	1992	Everlong The Colour And The Shape (Special Edi...

832 rows × 9 columns

- The final two queries shows what users were listening to the song 'Learn To Fly' from the artist'Foo Fighters' on their release 'There Is Nothing Left To Lose'. From both queries there are several instances where there were multiple playcounts for the song 'Learn To Fly' and 'The Pretender' which was one of the top recommended songs, so this could be an indicator that the recommendation system built here is working.

Observations and Insights:

- It looks like there is some similarities in that the first two songs in the recommendation list are songs from the same band and one song has the word 'fly' in it like the song we are looking up recommendations for based off the similarities. That is a pretty solid indication that the text features extraction worked to find similar songs. There are plenty songs in the database that have similar qualities it looks like so this would make for a more in-depth recommendation system for finding similar songs for a user.